

JIAQI GU

Graduate Research Assistant ◊ ECE Department ◊ University of Texas at Austin
jqgu@utexas.edu ◊ jeremielomelo.github.io

RESEARCH INTERESTS

Machine learning, optimization, efficient algorithm and architecture design for high-performance AI, software/hardware co-design, efficient AI computing with emerging technology, parallel computing/GPU acceleration for VLSI design automation

EDUCATION

University of Texas at Austin, TX, USA

Aug. 2018 – Present

Ph.D. student, Department of Electrical and Computer Engineering

Advisor: David Z. Pan

Co-advisor: Ray T. Chen

(GPA 4.0/4.0)

Fudan University, Shanghai, China

Sep. 2014 – Jun. 2018

B.E., Department of Microelectronic Science and Engineering (Eminent Engineer Program)

(GPA: 3.91/4.0)

(Rank top 2/71)

AWARDS AND HONORS

Donald O. Pederson Best Paper Award

IEEE TCAD 2021

Cockrell School Graduate Student Fellowship

UT Austin 2021

First Place at ACM Student Research Competition Grand Finals

ACM 2021

Best Poster Award at NSF Workshop on Machine Learning Hardware

NSF Workshop 2020

First Place at ACM/SIGDA Student Research Competition

ACM/SIGDA 2020

7th Place at IWLS Contest on Machine Learning+Logic Synthesis

IWLS 2020

DAC Young Fellow

DAC 2020,2021

Best Paper Finalist (1 out of 6)

DAC 2020

Best Paper Award

ASP-DAC 2020

4th Place, System Design Contest on Low Power Object Detection

DAC-SDC 2019

First Prize Scholarship

Fudan University 2017–2018

Top 5, HUAWEI & FUTURELAB AI Contest (CV Group)

Fudan University 2018

Top 11%, IEEEExtreme Global Programming Competition

IEEE 2017

2nd & 3rd Prize, National Mathematical Contest in Modeling

Fudan University 2016–2017

EXPERIENCE

Meta, CA, USA

May 2021 – Dec 2021

Research Intern, Meta reality labs, FAST AI team

- High-performance vision transformer backbone design with efficiency optimization for dense prediction vision tasks

University of Texas at Austin, TX, USA

Jan. 2019 – Present

Graduate Research Assistant

- **Memory-Efficient NN Design:** Designed memory-efficient multi-level low-rank weight generation methodology with mixed-precision quantization to save 10-20 \times on-chip memory cost for emerging NN accelerators;
- **Efficient Neural Arch. Design, NN Structured Pruning:** Designed hardware-efficient frequency-domain photonic neural network architecture; achieved 3-4 \times area reduction by using block-circulant matrices and structured pruning; further improved area and power by 2 \times and 10 \times by joint optimization and fine-grained pruning
- **NN Quantization and Robustness:** Developed differentiable quantization-aware training scheme in the unitary manifold to enable robust optical neural networks with low-precision voltage controls; achieved better accuracy and robustness with limited control resolution and device-level variations

- **NN On-Chip/On-Device Learning:** Developed customized CUDA extension for ONN simulation acceleration; proposed efficient on-chip learning algorithm for optical neural networks with stochastic zeroth-order optimization; achieved 3-4 \times higher learning efficiency, 10 \times better scalability, and better robustness than previous methods; proposed subspace optimization framework for 1,000 \times more scalable ONN on-chip learning
- **NN Efficient Training Framework:** Collaborated on developing efficient training framework for reversible neural architectures via constrained optimization; our dynamic programming based scheduling achieves 5-20% speedup with comparable memory efficiency when training reversible NNs
- **Photonics Neural Chip Tape-out:** Worked on photonic neural chip tape-out for novel ONN architectures using Advanced Micro Foundry (AMF); collaborated on the full-stack schematic design, layout, validation, tape-out, and measurement of photonic neural chips using PyTorch, Lumerical toolkits, and Synopsys optodesigner
- **GPU-Accelerated VLSI Detailed Placement:** Collaborated on developing GPU-accelerated concurrent VLSI detailed placement with CUDA; implemented and optimized global swap and parallel auction algorithm for batch-based independent-set-matching; achieved >10 \times speedup without quality drop
- **GPU-Accelerated VLSI Global Placement:** Collaborated on high-performance VLSI analytical global placement acceleration with CUDA on GPUs; optimized wirelength and density operators with CUDA; developed parallel congestion map estimation for routability optimization; achieved 40 \times speedup in global placement
- **VLSI Global Placement Algorithm:** Developed multi-electrostatics-based robust VLSI placement framework DREAMPlace 3.0 with PyTorch/C++/CUDA; proposed multi-electrostatic system for optimization under fence region constraints; developed divergence-aware optimizer for robust nonlinear global placement; achieved >13% HPWL improvement and >11% top5 overflow reduction compared with ISPD2015 contest winners
- **Efficient NN Learning and Power Optimization:** Proposed efficient ONN on-chip learning framework with two-level sparse optimizer and efficient power-aware optimization; achieved high convergence stability, \sim 10 \times training efficiency improvement, and \sim 10 \times power reduction than prior methods

University of Texas at Austin, TX, USA

Sep. 2018 – Jan. 2019

Graduate Research Assistant

- **FPGA Emulation of RISC-V Core:** Projected RISC-V Rocket Core on Zynq FPGA with Chisel3 and achieved communication between them
- **Fault Injection:** Customized FIRRTL transformation and built infrastructure for fault injection and system state snapshot

Fudan University, Shanghai, China

Aug. 2017 – Jul. 2018

Undergraduate Research Assistant

- **Medical Imaging Dataset:** Modified infant brain atlas and created complete tissue probability maps
- **MRI Reconstruction:** Developed two-stage reconstruction framework for infant thin-section MR image reconstruction by using GANs and CNNs; research is developing brand new method to improve reconstruction performance by fusing multi-planar MR images, and improving PSNR, SSIM, and NMI by 26.2%, 93.4%, and 25.3% respectively compared to bicubic interpolation
- **Ultra-sonic Image Processing:** Collaborated on super-resolution reconstruction of ultra-sonic imaging using U-Net and GANs; improved the full width at half maximum (FWHM) of point targets by 3.23%

PUBLICATIONS

Journal Papers

- [J10] Chenghao Feng*, **Jiaqi Gu***, Hanqing Zhu, Zhoufeng Ying, Zheng Zhao, David Z. Pan, and Ray T. Chen, “[Silicon photonic subspace neural chip for hardware-efficient deep learning](#),” *arXiv preprint 2111.06705*, 2021.
- [J9] Chenghao Feng, Zhoufeng Ying, Zheng Zhao, **Jiaqi Gu**, David Z. Pan, and Ray T. Chen, “[Towards high-speed and energy-efficient computing: A WDM-based scalable on-chip silicon integrated optical comparator](#),” *Laser & Photonics Reviews*, Jun. 2021.
- [J8] Zhoufeng Ying, Chenghao Feng, Zheng Zhao, **Jiaqi Gu**, Richard Soref, David Z. Pan, and Ray T. Chen, “[Sequential logic and pipelining in chip-based electronic-photonic digital computing](#),” *IEEE Photonics Journal*, Oct. 2020.
- [J7] **Jiaqi Gu**, Zheng Zhao, Chenghao Feng, Zhoufeng Ying, Mingjie Liu, Ray T. Chen, and David Z. Pan, “[Towards Hardware-Efficient Optical Neural Networks: Beyond FFT Architecture via Joint Learnability](#),” *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems (TCAD)*, 2020.

- [J6] Chenghao Feng, Zhoufeng Ying, Zheng Zhao, **Jiaqi Gu**, Ray T. Chen, and David Z. Pan, “Wavelength-division-multiplexing (WDM)-based integrated electronic–photonic switching network (EPSN) for high-speed data processing and transportation,” *Nanophotonics*, Aug. 2020.
- [J5] Yibo Lin, Zixuan Jiang, **Jiaqi Gu**, Wuxi Li, Shounak Dhar, Haoxing Ren, Brucek Khailany, and David Z. Pan, “DREAMPlace: Deep Learning Toolkit-Enabled GPU Acceleration for Modern VLSI Placement,” *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems (TCAD)*, Jun. 2020. (**Best Paper Award**)
- [J4] Zhoufeng Ying, Chenghao Feng, Zheng Zhao, Shounak Dhar, Hamed Dalir, **Jiaqi Gu**, Yue Cheng, Richard Soref, David Z. Pan, and Ray T. Chen, “Electronic-photonic Arithmetic Logic Unit for High-speed Computing,” *Nature Communications*, Apr. 2020.
- [J3] Yibo Lin, Wuxi Li, **Jiaqi Gu**, Mark Ren, Brucek Khailany, and David Z. Pan, “ABCDPlace: Accelerated Batch-based Concurrent Detailed Placement on Multi-threaded CPUs and GPUs,” *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems (TCAD)*, Feb. 2020.
- [J2] Ruoyao Wang, Zhenghan Fang, **Jiaqi Gu**, Yi Guo, Shicong Zhou, Yuanyuan Wang, Cai Chang, and Jinhua Yu, “High-resolution Image Reconstruction for Portable Ultrasound Imaging Devices,” *EURASIP Journal on Advances in Signal Processing*, Dec. 2019.
- [J1] **Jiaqi Gu**, Zeju Li, Yuanyuan Wang, Haowei Yang, Zhongwei Qiao, and Jinhua Yu, “Deep Generative Adversarial Networks for Thin-section Infant MR Image Reconstruction,” *IEEE Access*, May 2019.

Conference Papers

- [C34] Chenghao Feng, **Jiaqi Gu**, Hanqing Zhu, Zhoufeng Ying, Zheng Zhao, David Z. Pan, and Ray T. Chen, “Optoelectronically Interconnected Hardware-Efficient Deep Learning using Silicon Photonic Chips,” *Conference on Lasers and Electro-Optics*, Mar. 2022.
- [C33] Chenghao Feng, **Jiaqi Gu**, Hanqing Zhu, David Z. Pan, and Ray T. Chen, “Design and Experimental Demonstration of A Hardware-Efficient Integrated Optical Neural Network,” *Conference on Lasers and Electro-Optics*, Mar. 2022.
- [C32] **Jiaqi Gu**, Hyoukjun Kwon, Dilin Wang, Wei Ye, Meng Li, Yu-Hsin Chen, Liangzhen Lai, Vikas Chandra, and David Z. Pan, “Multi-Scale High-Resolution Vision Transformer for Semantic Segmentation,” *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022.
- [C31] **Jiaqi Gu**, Hanqing Zhu, Chenghao Feng, Zixuan Jiang, Mingjie Liu, Shuhan Zhang, Ray T. Chen, and David Z. Pan, “ADEPT: Automatic Differentiable DEsign of Photonic Tensor Cores,” *ACM/IEEE Design Automation Conference (DAC)*, Jul. 2022.
- [C30] Hanrui Wang, Zirui Li, **Jiaqi Gu**, Yongshan Ding, David Z. Pan, and Song Han, “On-Chip QNN: Towards Efficient On-Chip Training of Quantum Neural Networks,” *ACM/IEEE Design Automation Conference (DAC)*, Jul. 2022.
- [C29] Hanrui Wang, **Jiaqi Gu**, Yongshan Ding, Zirui Li, Frederic T. Chong, David Z. Pan, and Song Han, “RobustQNN: Noise-Aware Training for Robust Quantum Neural Networks,” *ACM/IEEE Design Automation Conference (DAC)*, Jul. 2022.
- [C28] Zizheng Guo, Mingjie Liu, **Jiaqi Gu**, Shuhan Zhang, David Z. Pan, and Yibo Lin, “A Timing Engine Inspired Graph Neural Network Model for Pre-Routing Slack Prediction,” *ACM/IEEE Design Automation Conference (DAC)*, Jul. 2022.
- [C27] Hanrui Wang, Yongshan Ding, **Jiaqi Gu**, Yujun Lin, David Z. Pan, Frederic T. Chong, and Song Han, “QuantumNAS: Noise-Adaptive Search for Robust Quantum Circuits,” *IEEE International Symposium on High Performance Computer Architecture (HPCA)*, Feb. 2022.
- [C26] Hanqing Zhu, **Jiaqi Gu**, Chenghao Feng, Mingjie Liu, Zixuan Jiang, Ray T. Chen, and David Z. Pan, “ELight: Enabling Efficient Photonic In-Memory Neurocomputing with Life Enhancement,” *IEEE/ACM Asia and South Pacific Design Automation Conference (ASPDAC)*, Jan. 2022.

- [C25] **Jiaqi Gu**, Hanqing Zhu, Chenghao Feng, Zixuan Jiang, Ray T. Chen, and David Z. Pan, “[L2ight: Enabling On-Chip Learning for Optical Neural Networks via Efficient in-situ Subspace Optimization](#),” *Conference on Neural Information Processing Systems (NeurIPS)*, Dec. 2021.
- [C24] **Jiaqi Gu**, Hanqing Zhu, Chenghao Feng, Mingjie Liu, Zixuan Jiang, Ray T. Chen, and David Z. Pan, “[Towards Memory-Efficient Neural Networks via Multi-Level in situ Generation](#),” *International Conference on Computer Vision (ICCV)*, Oct. 2021.
- [C23] Zixuan Jiang, **Jiaqi Gu**, and David Z. Pan, “[A New Acceleration Paradigm for Discrete Cosine Transform and Other Fourier-Related Transforms](#),” *arXiv preprint 2110.01172*, 2021.
- [C22] Zixuan Jiang, **Jiaqi Gu**, Mingjie Liu, Keren Zhu, and David Z. Pan, “[Optimizer Fusion: Efficient Training with Better Locality and Parallelism](#),” *International Conference on Learning Representations (ICLR) Workshop, Hardware Aware Efficient Training (HAET)*, May 2021.
- [C21] Chenghao Feng, **Jiaqi Gu**, Hanqing Zhu, David Z. Pan, and Ray T. Chen, “Experimental Demonstration of a WDM-based Integrated Optical Decoder for Compact Optical Computing,” *Conference on Lasers and Electro-Optics*, May 2021.
- [C20] Jason Midkiff, Ali Rostamian, Kyoung Min Yoo, Aref Asghari, Chao Wang, Chenghao Feng, Zhoufeng Ying, **Jiaqi Gu**, Haixia Mei, Ching-Wen Chang, James Fang, Alan Huang, Jong-Dug Shin, Xiaochuan Xu, Michael Bukshtab, David Z. Pan, and Ray T. Chen, “[Integrated Photonics for Computing, Interconnects and Sensing](#),” *Conference on Lasers and Electro-Optics*, May 2021. (**Invited Paper**)
- [C19] **Jiaqi Gu**, Chenghao Feng, Zheng Zhao, Zhoufeng Ying, Ray T. Chen, and David Z. Pan, “[Efficient On-Chip Learning for Optical Neural Networks Through Power-Aware Sparse Zeroth-Order Optimization](#),” *Association for the Advancement of Artificial Intelligence (AAAI)*, Feb. 2021.
- [C18] Shubham Rai, Walter Lau Neto, Yukio Miyasaka, Xinpei Zhang, Mingfei Yu, Qingyang Yi, Masahiro Fujita, Guilherme B. Manske, Matheus F. Pontes, Leomar S. da Rosa Junior, Marilton S. de Aguiar, Paulo F. Butzen, Po-Chun Chien, Yu-Shan Huang, Hoa-Ren Wang, Jie-Hong R. Jiang, **Jiaqi Gu**, Zheng Zhao, Zixuan Jiang, David Z. Pan, *et al.*, “[Logic Synthesis Meets Machine Learning: Trading Exactness for Generalization](#),” *IEEE/ACM Proceedings Design, Automation and Test in Europe (DATE)*, Feb. 2021.
- [C17] **Jiaqi Gu**, Chenghao Feng, Zheng Zhao, Zhoufeng Ying, Mingjie Liu, Ray T. Chen, and David Z. Pan, “SqueezeLight: Towards Scalable Optical Neural Networks with Multi-Operand Ring Resonators,” *IEEE/ACM Proceedings Design, Automation and Test in Europe (DATE)*, Feb. 2021.
- [C16] **Jiaqi Gu**, Zheng Zhao, Chenghao Feng, Zhoufeng Ying, Ray T. Chen, and David Z. Pan, “O2NN: Optical Neural Networks with Differential Detection-Enabled Optical Operands,” *IEEE/ACM Proceedings Design, Automation and Test in Europe (DATE)*, Feb. 2021.
- [C15] Chenghao Feng, **Jiaqi Gu**, Zhoufeng Ying, Zheng Zhao, Ray T. Chen, and David Z. Pan, “Scalable fast-Fourier-transform-based (FFT-based) integrated optical neural network for compact and energy-efficient deep learning,” *SPIE Photonics West*, Mar. 2021.
- [C14] Chenghao Feng, Zhoufeng Ying, Zheng Zhao, **Jiaqi Gu**, Ray T. Chen, and David Z. Pan, “Wavelength-division-multiplexing (WDM)-based integrated electronic–photonic switching network (EPSN) for high-speed data processing and transportation,” *SPIE Photonics West*, Mar. 2021.
- [C13] **Jiaqi Gu**, Zixuan Jiang, and David Z. Pan, “[DREAMPlace 3.0: Multi-Electrostatics Based Robustness VLSI Placement with Region Constraints](#),” *IEEE/ACM International Conference on Computer-Aided Design (ICCAD)*, Nov. 2020.
- [C12] Zixuan Jiang, Keren Zhu, Mingjie Liu, **Jiaqi Gu**, and David Z. Pan, “[An Efficient Training Framework for Reversible Neural Architectures](#),” *European Conference on Computer Vision (ECCV)*, Aug. 2020.
- [C11] **Jiaqi Gu**, Zheng Zhao, Chenghao Feng, Wuxi Li, Ray T. Chen, and David Z. Pan, “[FLOPS: Efficient On-Chip Learning for Optical Neural Networks Through Stochastic Zeroth-Order Optimization](#),” *ACM/IEEE Design Automation Conference (DAC)*, Jul. 2020. (**Best Paper Candidate**)

- [C10] Mario Miscuglio, Zibo Hu, Shurui Li, **Jiaqi Gu**, Aydin Babakhani, Puneet Gupta, Chee-Wei Wong, David Pan, Seth Bank, Hamed Dalir, and Volker J. Sorger, “[Massive parallelism Fourier-optic convolutional processor](#),” *Signal Processing in Photonic Communications (SPPCom)*, Jul. 2020.
- [C9] Mario Miscuglio, Zibo Hu, Shurui Li, **Jiaqi Gu**, Aydin Babakhani, Puneet Gupta, Chee-Wei Wong, David Z. Pan, Seth Bank, Hamed Dalir, and Volker J. Sorger, “[Million-channel parallelism Fourier-optic convolutional filter and neural network processor](#),” *Conference on Lasers and Electro-Optics*, May 2020.
- [C8] Chenghao Feng, Zhoufeng Ying, Zheng Zhao, **Jiaqi Gu**, Ray T. Chen, and David Z. Pan, “[Integrated WDM-based Optical Comparator for High-speed Computing](#),” *Conference on Lasers and Electro-Optics*, May 2020.
- [C7] Chenghao Feng, Zheng Zhao, Zhoufeng Ying, **Jiaqi Gu**, David Z. Pan, and Ray T. Chen, “[Compact design of On-chip Elman Optical Recurrent Neural Network](#),” *Conference on Lasers and Electro-Optics*, May 2020.
- [C6] **Jiaqi Gu**, Zheng Zhao, Chenghao Feng, Hanqing Zhu, Ray T. Chen, and David Z. Pan, “[ROQ: A Noise-Aware Quantization Scheme Towards Robust Optical Neural Networks with Low-bit Controls](#),” *IEEE/ACM Proceedings Design, Automation and Test in Europe (DATE)*, Mar. 2020.
- [C5] Mingjie Liu, Keren Zhu, **Jiaqi Gu**, Linxiao Shen, Xiyuan Tang, Nan Sun, and David Z. Pan, “[Towards Decrypting the Art of Analog Layout: Placement Quality Prediction via Transfer Learning](#),” *IEEE/ACM Proceedings Design, Automation and Test in Europe (DATE)*, Mar. 2020.
- [C4] Chenghao Feng, Zhoufeng Ying, Zheng Zhao, **Jiaqi Gu**, Ray T. Chen, and David Z. Pan, “[Wavelength-division-multiplexing-based electronic-photonic network for high-speed computing](#),” *SPIE, Smart Photonic and Optoelectronic Integrated Circuits XXII*, Feb. 2020.
- [C3] **Jiaqi Gu**, Zheng Zhao, Chenghao Feng, Mingjie Liu, Ray T. Chen, and David Z. Pan, “[Towards Area-Efficient Optical Neural Networks: An FFT-based Architecture](#),” *IEEE/ACM Asia and South Pacific Design Automation Conference (ASPDAC)*, Jan. 2020. (**Best Paper Award**)
- [C2] Zheng Zhao, **Jiaqi Gu**, Zhoufeng Ying, Chenghao Feng, Ray T. Chen, and David Z. Pan, “[Design Technology for Scalable and Robust Photonic Integrated Circuits](#),” *IEEE/ACM International Conference on Computer-Aided Design (ICCAD)*, 2019. (**Invited Paper**)
- [C1] **Jiaqi Gu**, Ruoyao Wang, Jian Wang, Jinmei Lai, and Qinghua Duan, “[Remote Embedded Simulation System for SW/HW Co-design Based On Dynamic Partial Reconfiguration](#),” *International Conference on ASIC (ASICON)*, 2017.

RELATED COURSES

- | | |
|---|--|
| • EE382N: Computer Architecture | <i>Prof. Dam Sunwoo</i> |
| • EE382N: High-Speed Computer Arithmetic I | <i>Prof. Earl Swartzlander</i> |
| • EE382N: Computer Architecture: Parallelism/Locality | <i>Prof. Mattan Erez</i> |
| • CS395T: Parallel Algorithm Scientific Computing | <i>Prof. George Biros</i> |
| • CS394R: Reinforcement Learning: Theory and Practice | <i>Prof. Peter Stone and Prof. Scott Niekum</i> |
| • EE382M: VLSI I | <i>Prof. Jacob A. Abraham</i> |
| • EE382M: VLSI Physical Design Automation | <i>Prof. David Z. Pan</i> |
| • EE382V: Cross-layer ML Alg./HW Co-design | <i>Prof. Mattan Erez and Prof. Michael Orshansky</i> |
| • EE382M: VLSI CAD and Optimization | <i>Prof. David Z. Pan</i> |
| • EE381V: Combinatorial Optimization | <i>Prof. Constantine Caramanis</i> |
| • EE381V: Advanced Topics in Computer Vision | <i>Prof. Zhangyang (Atlas) Wang</i> |

SKILLS

Programming Languages

Python (PyTorch/TensorFlow), C/C++, CUDA, Matlab, Verilog

EDA Tools

Cadence Virtuoso, Synopsys Design Compiler, Hspice, Xilinx Vivado Design Suite, Synopsys Optodesigner