

JIAQI GU

Graduate Research Assistant ◊ ECE Department ◊ University of Texas at Austin
jqgu@utexas.edu ◊ <https://jeremiemelo.github.io>

RESEARCH INTERESTS

Machine learning, algorithm and architecture design, software/hardware co-design for neuromorphic computing, and GPU acceleration for VLSI physical design automation

EDUCATION

University of Texas at Austin, TX, USA

Aug. 2018 – Present

Ph.D. student, Department of Electrical and Computer Engineering

Advisor: David Z. Pan

Co-advisor: Ray T. Chen

(GPA 4.0/4.0)

Fudan University, Shanghai, China

Sep. 2014 – Jun. 2018

B.E., Department of Microelectronic Science and Engineering (Eminent Engineer Program)

(GPA: 3.91/4.0)

(Rank top 2/71)

EXPERIENCE

University of Texas at Austin, TX, USA

Jan. 2019 – Present

Graduate Research Assistant

- Designed novel frequency-domain photonic neuromorphic computing architecture for area-efficient optical neural network; achieved 3-4 \times area reduction by using block-circulant matrices and structured pruning compared with previous ONN architectures
- Developed differentiable quantization-aware training scheme in the unitary manifold to enable robust optical neural networks with low-precision voltage controls; achieved better accuracy and robustness with limited control resolution and device-level variations
- Proposed efficient ONN on-chip learning algorithm for optical neural networks with stochastic zeroth-order optimization algorithms; achieved 3-4 \times higher learning efficiency, 10 \times better scalability, and better robustness to thermal variations than previous methods
- Worked on photonic chip tapeout for novel ONN architectures with AMF PDKs using Synopsys optodesigner and Lumerical toolkits
- Collaborated on the full-stack schematic design, layout, validation, tape-out, and measurement of photonic neural chips using PyTorch, Lumerical toolkits, and Synopsys optodesigner
- Collaborated on developing GPU-accelerated concurrent VLSI detailed placement with CUDA; implemented and optimized global swap and parallel auction algorithm for batched-based independent-set-matching; achieved >10 \times speedup than sequential implementations without quality degradation
- Collaborated on high-performance VLSI analytical global placement acceleration with CUDA on GPUs; optimized wirelength and density computation operators with CUDA; developed parallel RUDY/RISA congestion map for routability optimization; achieved 40 \times speedup in global placement
- Developed multi-electrostatics-based robust VLSI placement framework DREAMPlace 3.0 with PyTorch/C++/CUDA; proposed multielectrostatic system for optimization under fence region constraints; developed divergence-aware optimizer for robust nonlinear global placement; achieved >13% HPWL improvement and >11% top5 overflow reduction compared with ISPD2015 contest winners
- Proposed efficient ONN on-chip learning framework with power-aware mixed-training strategy and two-level algorithmic sparsity; achieved much better training stability, \sim 10 \times training efficiency improvement, and \sim 10 \times power reduction than prior methods

University of Texas at Austin, TX, USA

Sep. 2018 – Jan. 2019

Graduate Research Assistant

- Projected RISC-V Rocket Core on Zynq FPGA with Chisel3 and achieved communication between them
- Customized FIRRTL transformation and built infrastructure for fault injection and system state snapshot

- Modified infant brain atlas offered by UNC and created complete tissue probability maps
- Developed two-stage reconstruction framework for infant thin-section MR image reconstruction by using GANs and CNN; research is developing brand new method to improve reconstruction performance by fusing multi-planar MR images, and improving PSNR, SSIM, and NMI by 26.2%, 93.4%, and 25.3% respectively compared to bicubic interpolation
- Collaborated on super-resolution reconstruction of ultra-sonic imaging using U-Net and GANs; improved the full width at half maximum (FWHM) of point targets by 3.23%

PUBLICATIONS

Journal Papers

- [J1] Zhoufeng Ying, Chenghao Feng, Zheng Zhao, **Jiaqi Gu**, Richard Soref, David Z. Pan, and Ray T. Chen, “[Sequential logic and pipelining in chip-based electronic-photonic digital computing](#),” *IEEE Photonics Journal*, Oct. 2020.
- [J2] **Jiaqi Gu**, Zheng Zhao, Chenghao Feng, Zhoufeng Ying, Mingjie Liu, Ray T. Chen, and David Z. Pan, “[Towards Hardware-Efficient Optical Neural Networks: Beyond FFT Architecture via Joint Learnability](#),” *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems (TCAD)*, 2020.
- [J3] Chenghao Feng, Zhoufeng Ying, Zheng Zhao, **Jiaqi Gu**, Ray T. Chen, and David Z. Pan, “[Wavelength-division-multiplexing \(WDM\)-based integrated electronic-photonic switching network \(EPSN\) for high-speed data processing and transportation](#),” *Nanophotonics*, Aug. 2020.
- [J4] Yibo Lin, Zixuan Jiang, **Jiaqi Gu**, Wuxi Li, Shounak Dhar, Haoxing Ren, Brucek Khailany, and David Z. Pan, “[DREAMPlace: Deep Learning Toolkit-Enabled GPU Acceleration for Modern VLSI Placement](#),” *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems (TCAD)*, Jun. 2020.
- [J5] Zhoufeng Ying, Chenghao Feng, Zheng Zhao, Shounak Dhar, Jiaqi Gu Hamed Dalir, **Yue Cheng**, Richard Soref, David Z. Pan, and Ray T. Chen, “[Electronic-photonic Arithmetic Logic Unit for High-speed Computing](#),” *Nature Communications*, Apr. 2020.
- [J6] Yibo Lin, Wuxi Li, **Jiaqi Gu**, Mark Ren, Brucek Khailany, and David Z. Pan, “[ABCDPlace: Accelerated Batch-based Concurrent Detailed Placement on Multi-threaded CPUs and GPUs](#),” *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems (TCAD)*, Feb. 2020.
- [J7] Ruoyao Wang, Zhenghan Fang, **Jiaqi Gu**, Yi Guo, Shicong Zhou, Yuanyuan Wang, Cai Chang, and Jinhua Yu, “[High-resolution Image Reconstruction for Portable Ultrasound Imaging Devices](#),” *EURASIP Journal on Advances in Signal Processing*, Dec. 2019.
- [J8] **Jiaqi Gu**, Zeju Li, Yuanyuan Wang, Haowei Yang, Zhongwei Qiao, and Jinhua Yu, “[Deep Generative Adversarial Networks for Thin-section Infant MR Image Reconstruction](#),” *IEEE Access*, May 2019.

Conference Papers

- [C1] **Jiaqi Gu**, Chenghao Feng, Zheng Zhao, Zhoufeng Ying, Ray T. Chen, and David Z. Pan, “[Efficient On-Chip Learning for Optical Neural Networks Through Power-Aware Sparse Zeroth-Order Optimization](#),” *Association for the Advancement of Artificial Intelligence (AAAI)*, Feb. 2021. (accepted)
- [C2] Shubham Rai, Walter Lau Neto, Yukio Miyasaka, Xinpei Zhang, Mingfei Yu, Qingyang Yi, Masahiro Fujita, Guilherme B. Manske, Matheus F. Pontes, Leomar S. da Rosa Junior, Marilton S. de Aguiar, Paulo F. Butzen, Po-Chun Chien, Yu-Shan Huang, Hoa-Ren Wang, Jie-Hong R. Jiang, **Jiaqi Gu**, Zheng Zhao, Zixuan Jiang, David Z. Pan, *et al.*, “[Logic Synthesis Meets Machine Learning: Trading Exactness for Generalization](#),” *IEEE/ACM Proceedings Design, Automation and Test in Europe (DATE)*, Feb. 2021. (accepted)

- [C3] **Jiaqi Gu**, Chenghao Feng, Zheng Zhao, Zhoufeng Ying, Mingjie Liu, Ray T. Chen, and David Z. Pan, “SqueezeLight: Towards Scalable Optical Neural Networks with Multi-Operand Ring Resonators,” *IEEE/ACM Proceedings Design, Automation and Test in Europe (DATE)*, Feb. 2021. (accepted)
- [C4] **Jiaqi Gu**, Zheng Zhao, Chenghao Feng, Zhoufeng Ying, Ray T. Chen, and David Z. Pan, “O2NN: Optical Neural Networks with Differential Detection-Enabled Optical Operands,” *IEEE/ACM Proceedings Design, Automation and Test in Europe (DATE)*, Feb. 2021. (accepted)
- [C5] Chenghao Feng, **Jiaqi Gu**, Zhoufeng Ying, Zheng Zhao, Ray T. Chen, and David Z. Pan, “Scalable fast-Fourier-transform-based (FFT-based) integrated optical neural network for compact and energy-efficient deep learning,” *SPIE Photonics West*, Mar. 2021. (accepted)
- [C6] **Jiaqi Gu**, Zixuan Jiang, and David Z. Pan, “[DREAMPlace 3.0: Multi-Electrostatics Based Robustness VLSI Placement with Region Constraints](#),” *IEEE/ACM International Conference on Computer-Aided Design (ICCAD)*, Nov. 2020.
- [C7] Zixuan Jiang, Keren Zhu, Mingjie Liu, **Jiaqi Gu**, and David Z. Pan, “[An Efficient Training Framework for Reversible Neural Architectures](#),” *European Conference on Computer Vision (ECCV)*, Aug. 2020.
- [C8] **Jiaqi Gu**, Zheng Zhao, Chenghao Feng, Wuxi Li, Ray T. Chen, and David Z. Pan, “[FLOPS: Efficient On-Chip Learning for Optical Neural Networks Through Stochastic Zeroth-Order Optimization](#),” *ACM/IEEE Design Automation Conference (DAC)*, Jul. 2020. (**Best Paper Candidate**)
- [C9] Mario Miscuglio, Zibo Hu, Shurui Li, **Jiaqi Gu**, Aydin Babakhani, Puneet Gupta, Chee-Wei Wong, Hamed Dalir David Z. Pan Seth Bank, and Volker J. Sorger, “[Million-channel parallelism Fourier-optic convolutional filter and neural network processor](#),” *Conference on Lasers and Electro-Optics*, May 2020.
- [C10] Chenghao Feng, Zhoufeng Ying, Zheng Zhao, **Jiaqi Gu**, Ray T. Chen, and David Z. Pan, “[Integrated WDM-based Optical Comparator for High-speed Computing](#),” *Conference on Lasers and Electro-Optics*, May 2020.
- [C11] Chenghao Feng, Zheng Zhao, Zhoufeng Ying, **Jiaqi Gu**, David Z. Pan, and Ray T. Chen, “[Compact design of On-chip Elman Optical Recurrent Neural Network](#),” *Conference on Lasers and Electro-Optics*, May 2020.
- [C12] **Jiaqi Gu**, Zheng Zhao, Chenghao Feng, Hanqing Zhu, Ray T. Chen, and David Z. Pan, “[ROQ: A Noise-Aware Quantization Scheme Towards Robust Optical Neural Networks with Low-bit Controls](#),” *IEEE/ACM Proceedings Design, Automation and Test in Europe (DATE)*, Mar. 2020.
- [C13] Mingjie Liu, Keren Zhu, **Jiaqi Gu**, Linxiao Shen, Xiyuan Tang, Nan Sun, and David Z. Pan, “[Towards Decrypting the Art of Analog Layout: Placement Quality Prediction via Transfer Learning](#),” *IEEE/ACM Proceedings Design, Automation and Test in Europe (DATE)*, Mar. 2020.
- [C14] Chenghao Feng, Zhoufeng Ying, **Zheng Zhao**, Jiaqi Gu, Ray T. Chen, and David Z. Pan, “[Wavelength-division-multiplexing-based electronic-photonic network for high-speed computing](#),” *SPIE, Smart Photonic and Optoelectronic Integrated Circuits XXII*, Feb. 2020.
- [C15] **Jiaqi Gu**, Zheng Zhao, Chenghao Feng, Mingjie Liu, Ray T. Chen, and David Z. Pan, “[Towards Area-Efficient Optical Neural Networks: an FFT-based architecture](#),” *IEEE/ACM Asia and South Pacific Design Automation Conference (ASPDAC)*, Jan. 2020. (**Best Paper Award**)
- [C16] Zheng Zhao, **Jiaqi Gu**, Zhoufeng Ying, Chenghao Feng, Ray T. Chen, and David Z. Pan, “[Design Technology for Scalable and Robust Photonic Integrated Circuits](#),” *IEEE/ACM International Conference on Computer-Aided Design (ICCAD)*, 2019. (**Invited Paper**)
- [C17] **Jiaqi Gu**, Ruoyao Wang, Jian Wang, Jinmei Lai, and Qinghua Duan, “[Remote Embedded Simulation System for SW/HW Co-design Based On Dynamic Partial Reconfiguration](#),” *International Conference on ASIC (ASICON)*, 2017.

RELATED COURSES

- EE382N: Computer Architecture
- EE382N: High-Speed Computer Arithmetic I

Prof. Dam Sunwoo

Prof. Earl Swartzlander

- EE382N: Computer Architecture: Parallelism/Locality *Prof. Mattan Erez*
- CS395T: Parallel Algorithm Scientific Computing *Prof. George Biros*
- CS394R: Reinforcement Learning: Theory and Practice *Prof. Peter Stone and Prof. Scott Niekum*
- EE382M: VLSI I *Prof. Jacob A. Abraham*
- EE382M: VLSI Physical Design Automation *Prof. David Z. Pan*
- EE382V: Cross-layer ML Alg./HW Co-design *Prof. Mattan Erez and Prof. Michael Orshansky*
- EE382M: VLSI CAD and Optimization *Prof. David Z. Pan*
- EE381V: Combinatorial Optimization *Prof. Constantine Caramanis*

SKILLS

Programming Languages

Python (PyTorch/TensorFlow), C/C++, CUDA, Matlab, Verilog

EDA Tools

Cadence Virtuoso, Synopsys Design Compiler, Hspice, Xilinx Vivado Design Suite, Synopsys Optodesigner

AWARDS AND HONORS

| | | |
|---|------------------|-----------|
| Best Poster Award at NSF Workshop on Machine Learning Hardware | NSF Workshop | 2020 |
| Gold Medal at ACM/SIGDA Student Research Competition | ACM/SIGDA | 2020 |
| 7th Place at IWLS Contest on Machine Learning+Logic Synthesis | IWLS | 2020 |
| DAC Young Fellow | DAC | 2020 |
| Best Paper Finalist (1 out of 6) | DAC | 2020 |
| Best Paper Award | ASP-DAC | 2020 |
| 4th Place, System Design Contest on Low Power Object Detection | DAC-SDC | 2019 |
| First Prize Scholarship | Fudan University | 2017–2018 |
| Top 5, HUAWEI & FUTURELAB AI Contest (CV Group) | Fudan University | 2018 |
| Top 11%, IEEEExtreme Global Programming Competition | IEEE | 2017 |
| 2nd & 3rd Prize, National Mathematical Contest in Modeling | Fudan University | 2016–2017 |